

$$p(\Theta|D) = \frac{p(D|\Theta)p(\Theta)}{p(D|\Theta)p(\Theta) + p(D|\neg\Theta)p(\neg\Theta)}$$

Bayesian Learning 732A46: Lecture 10

Matias Quiroz^{1,2}

¹Division of Statistics and Machine Learning, Linköping University

²Research Division, Sveriges Riksbank

May 2016

- ▶ Bayesian model comparison
- ▶ Computing marginal likelihoods
- ▶ Bayesian model averaging

Using the likelihood for model comparison

- ▶ Consider two models for the data $y = (y_1, \dots, y_n)$: M_1 and M_2 .
- ▶ Let $p_k(y|\theta_k)$ denote the **data density** (fixed θ_k) under model M_k .
- ▶ If we know θ_1 and θ_2 , the **likelihood ratio** is useful

$$\frac{p_1(y|\theta_1)}{p_2(y|\theta_2)}.$$

- ▶ But often we **do not know** θ_1 and θ_2 .
- ▶ **Frequentist**: The **likelihood ratio** with the **MLE** plugged in:

$$\frac{p_1(y|\hat{\theta}_1)}{p_2(y|\hat{\theta}_2)}.$$

- ▶ **Bigger models** always win with estimated likelihood ratio.
- ▶ **Hypothesis tests** become problematic for non-nested models.

Bayesian model comparison

- ▶ Use your priors $p_1(\theta_1)$ and $p_2(\theta_2)$ to get rid (**average over**) of θ .
- ▶ The **marginal likelihood** for model M_k with parameters θ_k

$$p_k(y) = \int p_k(y|\theta_k)p_k(\theta_k)d\theta_k.$$

- ▶ Recall **Bayes' theorem** in the simple case of $\theta = \{H, H^c\}$

$$\Pr(H|E) = \frac{\Pr(E|H)\Pr(H)}{\Pr(E)}, \quad \Pr(E) = \Pr(E|H)\Pr(H) + \Pr(E|H^c)\Pr(H^c)$$

The marginal likelihood in words

The **marginal likelihood** $\Pr(E)$ is a **weighted average** of the probability of the evidence under the different hypothesis. The weights **are given by the prior probabilities**.

- ▶ θ_k (or H, H^c) is removed (**averaged out**) by the prior. **Priors matter!**

- ▶ The **Bayes factor**

$$B_{12}(y) = \frac{p_1(y)}{p_2(y)}.$$

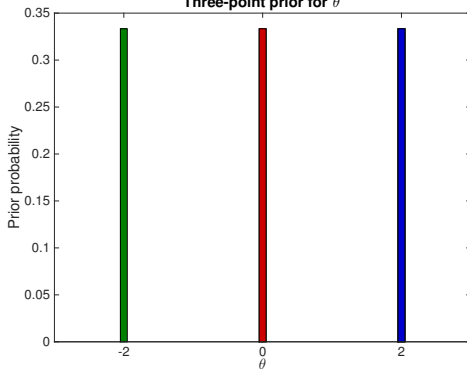
- ▶ **Bayesian machinery**: Posterior model probabilities

$$\underbrace{\Pr(M_k|y)}_{\text{Posterior model prob.}} \propto \underbrace{p(y|M_k)}_{\text{marginal likelihood [=} p_k(y)\text{]}} \cdot \underbrace{\Pr(M_k)}_{\text{prior model prob.}}$$

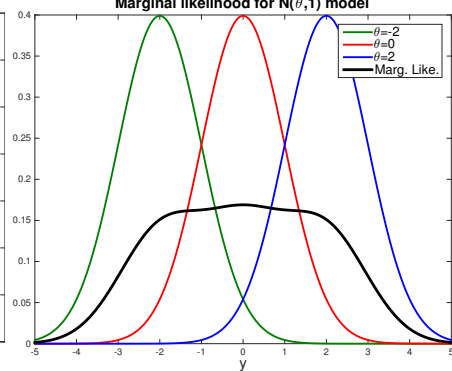
- ▶ **Important**: Two sets of priors
 1. Prior for **the parameters** θ_k within model M_k ("**the usual**" prior)
 2. Prior for **the models** $\Pr(M_k)$.

Priors matter

Three-point prior for θ



Marginal likelihood for $N(\theta, 1)$ model



Example: Geometric vs Poisson

- ▶ Model 1 - **Geometric** with **Beta** prior:

- ▶ $y_1, \dots, y_n | \theta_1 \sim \text{Geometric}(\theta_1),$

$$p(y_i | \theta_1) = (1 - \theta_1)^{y_i} \theta_1 \quad y_i \in \{0, 1, 2, \dots\}, 0 \leq \theta_1 \leq 1.$$

- ▶ $\theta_1 \sim \text{Beta}(\alpha_1, \beta_1),$

$$p(\theta_1) = \frac{\Gamma(\alpha_1 + \beta_1)}{\Gamma(\alpha_1)\Gamma(\beta_1)} \theta_1^{\alpha_1-1} (1 - \theta_1)^{\beta_1-1}.$$

- ▶ Model 2 - **Poisson** with **Gamma** prior:

- ▶ $y_1, \dots, y_n | \theta_2 \sim \text{Poisson}(\theta_2),$

$$p(y_i | \theta_2) = \frac{\theta_2^{y_i} \exp(-\theta_2)}{y_i!} \quad y_i \in \{0, 1, 2, \dots\}, \theta_2 > 0.$$

- ▶ $\theta_2 \sim \text{Gamma}(\alpha_2, \beta_2),$

$$p(\theta_2) = \frac{\beta_2^{\alpha_2}}{\Gamma(\alpha_2)} \theta_2^{\alpha_2-1} \exp(-\beta_2 \theta_2).$$

Geometric vs Poisson: $p(y)$ for Geometric (M_1)

- **Marginal likelihood** for M_1 [$y = (y_1, \dots, y_n)$]

$$\begin{aligned} p_1(y) &= \int p_1(y|\theta_1)p(\theta_1)d\theta_1 \\ &= \int \left(\prod_{i=1}^n p(y_i|\theta_1) \right) p(\theta_1)d\theta_1 \\ &= \frac{\Gamma(\alpha_1 + \beta_1)}{\Gamma(\alpha_1)\Gamma(\beta_1)} \int (1 - \theta_1)^{\sum_{i=1}^n y_i} \theta_1^n \times \theta_1^{\alpha_1-1} (1 - \theta_1)^{\beta_1-1} d\theta_1 \\ &= \frac{\Gamma(\alpha_1 + \beta_1)}{\Gamma(\alpha_1)\Gamma(\beta_1)} \int \theta_1^{n+\alpha_1-1} (1 - \theta_1)^{n\bar{y}+\beta_1-1} d\theta_1 \end{aligned}$$

- The **beta function**

$$B(a, b) = \int_0^1 t^{a-1} (1 - t)^{b-1}, \quad a, b > 0.$$

- **Nice property** of the beta function

$$B(a, b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}.$$

► Thus

$$\begin{aligned} p_1(y) &= \frac{\Gamma(\alpha_1 + \beta_1)}{\Gamma(\alpha_1)\Gamma(\beta_1)} \int \theta_1^{n+\alpha_1-1} (1 - \theta_1)^{n\bar{y}+\beta_1-1} d\theta_1 \\ &= \frac{\Gamma(\alpha_1 + \beta_1)}{\Gamma(\alpha_1)\Gamma(\beta_1)} B(n + \alpha_1, n\bar{y} + \beta_1) \\ &= \frac{\Gamma(\alpha_1 + \beta_1)}{\Gamma(\alpha_1)\Gamma(\beta_1)} \frac{\Gamma(n + \alpha_1)\Gamma(n\bar{y} + \beta_1)}{\Gamma(n + \alpha_1 + n\bar{y} + \beta_1)}. \end{aligned}$$

► **Note:** It **does not** depend on θ_1 . θ_1 has been averaged out!

Geometric vs Poisson: $p(y)$ for Poisson (M_2)

- **Marginal likelihood** for M_2 [$y = (y_1, \dots, y_n)$]

$$\begin{aligned} p_2(y) &= \int p_2(y|\theta_2) p(\theta_2) d\theta_2 \\ &= \int \left(\prod_{i=1}^n p(y_i|\theta_2) \right) p(\theta_2) d\theta_2 \\ &= \frac{\beta_2^{\alpha_2}}{\Gamma(\alpha_2)} \int \frac{\theta_2^{\sum y_i}}{\prod_{i=1}^n y_i} \exp(-n\theta_2) \times \theta_2^{\alpha_2-1} \exp(-\beta_2\theta_2) d\theta_2 \\ &= \frac{\beta_2^{\alpha_2}}{\Gamma(\alpha_2) \prod_{i=1}^n y_i} \int \theta_2^{n\bar{y} + \alpha_2 - 1} \exp(-(n + \beta_2)\theta_2) d\theta_2 \end{aligned}$$

- The **gamma function**

$$\Gamma(c) = \int_0^\infty t^{c-1} \exp(-t) dt, \quad c > 0.$$

- ... rewritten to fit **our form above** (simple change of variables) ...

$$\frac{1}{(n + \beta_2)^c} \Gamma(c) = \int_0^\infty t^{c-1} \exp(-(n + \beta_2)t) dt, \quad c > 0.$$

► Thus

$$\begin{aligned} p_2(y) &= \frac{\beta_2^{\alpha_2}}{\Gamma(\alpha_2) \prod_{i=1}^n y_i} \int \theta_2^{n\bar{y} + \alpha_2 - 1} \exp(-(n + \beta_2)\theta_2) d\theta_2 \\ &= \frac{\beta_2^{\alpha_2} \Gamma(n\bar{y} + \alpha_2)}{\Gamma(\alpha_2)(n + \beta_2)^{n\bar{y} + \alpha_2} \prod_{i=1}^n y_i}. \end{aligned}$$

► **Note (again!)**: It **does not** depend on θ_2 . θ_2 has been averaged out!

Geometric vs Poisson, cont.

- ▶ **Before** comparing the results we need to set the hyper-parameters in **some suitable way**.
- ▶ Set **hyper-parameters** so that the prior predictive means match

$$E(y_i|M_1) = E(y_i|M_2) \implies (\alpha_1 - 1)\alpha_2 = \beta_1\beta_2$$


- ▶ The **prior predictive mean** computed by the **tower property**

$$E(y_i|M_k) = E_{\theta} (E_{y_i|\theta}(y_i|\theta, M_k)) , \quad \text{for } k = 1, 2,$$


and

$$E_{y_i|\theta}(y_i|\theta, M_k) = \begin{cases} \frac{\theta_1}{1-\theta_1}, & \text{if } k = 1 \\ \theta_2, & \text{if } k = 2. \end{cases}$$

Geometric vs Poisson for Pois(1) data



Geometric vs Poisson for Pois(1) data



Properties of Bayesian model comparison

- ▶ **Coherence** of pair-wise comparisons

$$B_{12} = B_{13} \cdot B_{32}.$$

- ▶ **Consistency** when true model is in $\mathcal{M} = \{M_1, \dots, M_K\}$

$$\Pr(M = M_{TRUE}|y) \rightarrow 1 \quad \text{as } n \rightarrow \infty.$$

- ▶ **“KL-consistency”** when $M_{TRUE} \notin \mathcal{M}$

$$\Pr(M = M^*|y) \rightarrow 1 \quad \text{as } n \rightarrow \infty,$$

where M^* is the model that minimizes Kullback-Leibler distance

$$D_{KL}(p_{TRUE}, p_M) = \int p_{TRUE}(y) \log \left(\frac{p_M(y)}{p_{TRUE}(y)} \right) dy$$

between $p_M(y)$ and $p_{TRUE}(y)$.

Some warnings

- ▶ Smaller models **always win** when priors are very vague.
- ▶ **Improper priors can't be used** for model comparison.
- ▶ **Bayes factors** are **relative measures**! **Does not** say anything about a single model's adequacy.

Bayesian hypothesis testing

- ▶ **Hypothesis testing** is a **model selection** problem.

- ▶ **Example:** Bernoulli model with prior $\theta \sim \text{Beta}(\alpha, \beta)$

$$M_0 : y_1, \dots, y_n | \theta_0 \stackrel{iid}{\sim} \text{Bernoulli}(\theta_0)$$

$$M_1 : y_1, \dots, y_n | \theta \stackrel{iid}{\sim} \text{Bernoulli}(\theta).$$

- ▶ **Likelihood:** $p(y|\theta) = \theta^s(1 - \theta)^f$ ($y = (y_1, \dots, y_n)$, $s = \sum y_i$, $f = n - s$).

- ▶ **Marginal likelihoods**

- ▶ For model M_1

$$p(y|M_1) = \theta^s(1 - \theta)^f.$$

- ▶ For model M_2

$$\begin{aligned} p(y|M_2) &= \int \theta^s(1 - \theta)^f \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1}(1 - \theta)^{\beta-1} d\theta \\ &= \frac{\Gamma(\alpha + \beta)\Gamma(s + \alpha)\Gamma(f + \beta)}{\Gamma(\alpha)\Gamma(\beta)\Gamma(n + \alpha + \beta)}. \end{aligned}$$

- ▶ Reject (or accept) based on the **posterior model probabilities**


$$Pr(M_k|y) \propto p(y|M_k)Pr(M_k), \text{ for } k = 0, 1.$$

- ▶ A **"sharp null"** hypothesis is equivalent to using '**spike-and-slab**' prior:

$$p(\theta) = \pi\delta_{\theta_0}(\theta) + (1 - \pi)\text{Beta}(\alpha, \beta).$$

- ▶ Think about the **shrinkage mechanism**!
- ▶ **Note**: data can now **support** a null hypothesis (not only reject it).

Spike-and-slab prior [with $\theta_0 = 0.5$]



π

Marginal likelihood - a measure of out-of-sample predictive performance

- ▶ **The marginal likelihood** can be decomposed as

$$p(y_1, \dots, y_n) = p(y_1)p(y_2|y_1) \cdots p(y_n|y_1, y_2, \dots, y_{n-1}).$$

- ▶ Assume that y_i is **independent** of y_1, \dots, y_{i-1} **conditional** on θ :

$$p(y_i|y_1, \dots, y_{i-1}) = \int p(y_i|\theta)p(\theta|y_1, \dots, y_{i-1})d\theta$$

- ▶ **The prediction** of y_1 is **based on the prior** of θ , and is therefore **sensitive to the prior**.
- ▶ In contrast, **the prediction** of y_n **uses almost all the data** to infer θ . If n is large **influence of prior is negligible** for y_n .
- ▶ **Summary:** "Early" out-of-sample predictions are more influenced by $p(\theta)$.

Illustrating the sensitivity to the prior for early obs

- ▶ **Model:** $y_1, \dots, y_n | \theta \sim \mathcal{N}(\theta, \sigma^2)$ with σ^2 **known**.
- ▶ **Prior:** $\theta \sim \mathcal{N}(0, \kappa^2 \sigma^2)$ [for simplified expressions].
- ▶ **Partial posterior** up to observation $i - 1$ ($\mu_0 = 0$)

$$\theta | y_1, \dots, y_{i-1} \sim \mathcal{N} \left[w_i(\kappa) \cdot \bar{y}_{i-1}, \frac{\sigma^2}{i - 1 + \kappa^{-2}} \right]$$


where $w_i(\kappa) = \frac{i-1}{i-1+\kappa^{-2}}$ [the usual weighted average story].

- ▶ **Predictive density** for obs $i - 1$


$$y_i | y_1, \dots, y_{i-1} \sim \mathcal{N} \left[w_i(\kappa) \cdot \bar{y}_{i-1}, \sigma^2 \left(1 + \frac{1}{i - 1 + \kappa^{-2}} \right) \right].$$

- ▶ **Terms with i large:** $y_i | y_1, \dots, y_{i-1} \overset{\text{approx}}{\sim} \mathcal{N}(\bar{y}_{i-1}, \sigma^2)$, **not sensitive** to κ
- ▶ For $i = 1$, $y_1 \sim \mathcal{N} \left[0, \sigma^2 \left(1 + \frac{1}{\kappa^{-2}} \right) \right]$ can be **very sensitive** to κ .

First observation is sensitive to κ



First observation is sensitive to κ



Log Predictive Score - LPS: a way to reduce the sensitivity

- ▶ **Simple idea:** a measure similar to the marginal likelihood but where **the first observation is less sensitive** to the prior.
- ▶ **Sacrifice** n^* observations to train/update the prior.
- ▶ **Predictive density score:** PS

$$PS(n^*) = p(y_{n^*+1}|y_1, \dots, y_{n^*}) \cdots p(y_n|y_1, \dots, y_{n-1}).$$

- ▶ **Compare** PS to $p(y)$ in factorized form.
- ▶ Usually report on log scale: **Log Predictive Score (LPS)**.
- ▶ Which observations to **train/update** with (and which to predict)?
- ▶ **Split the data:** *Training* and *test* data
 - ▶ Straightforward for **time series**.
 - ▶ **Cross-sectional data:** cross-validation is useful.

Computing the marginal likelihood: Conjugate models

- ▶ Computing the **marginal likelihood** requires integration w.r.t. θ .
- ▶ **Short cut** for **conjugate models** by rearrangement of Bayes' theorem:

$$p(y) = \frac{p(y|\theta)p(\theta)}{p(\theta|y)}.$$

- ▶ By conjugacy $p(\theta|y)$ is **analytically available**.
- ▶ Insert everything and **work out the algebra**.

Computing the marginal likelihood: Simulation methods

- Usually difficult (or **impossible**) to analytically derive

$$p(y) = \int p(y|\theta)p(\theta)d\theta = E_{\theta}[p(y|\theta)].$$

- Draw from the prior $\theta^{(1)}, \dots, \theta^{(N)}$ and use the usual **Monte Carlo estimate**

$$\hat{p}(y) = \frac{1}{N} \sum_{i=1}^N p(y|\theta^{(i)}).$$

- **Unstable** (huge variance) if the likelihood is somewhat different from the prior.
- **Importance sampling**. Let $\theta^{(1)}, \dots, \theta^{(N)}$ be iid draws from $g(\theta)$.

$$\int p(y|\theta)p(\theta)d\theta = \int \frac{p(y|\theta)p(\theta)}{g(\theta)}g(\theta)d\theta \approx \frac{1}{N} \sum_{i=1}^N \frac{p(y|\theta^{(i)})p(\theta^{(i)})}{g(\theta^{(i)})}.$$

- **Modified Harmonic mean**: $g(\theta) = \mathcal{N}(\tilde{\theta}, \tilde{\Sigma}) \cdot I_c(\theta)$, where $\tilde{\theta}$ and $\tilde{\Sigma}$ is the posterior mean and covariance matrix estimated from an MCMC chain, and $I_c(\theta) = 1$ if $(\theta - \tilde{\theta})' \tilde{\Sigma}^{-1}(\theta - \tilde{\theta}) \leq c$.

Computing the marginal likelihood: Simulation methods, cont.

- ▶ Rearrangement of **Bayes' theorem** (again!): $p(y) = p(y|\theta)p(\theta)/p(\theta|y)$.
- ▶ **Note 1**: Need the full expression for the posterior, **including** the constants ind of θ .
- ▶ **Note 2**: LHS is **independent** of θ . RHS **depends** on θ ...
- ▶ ... any θ must cancel. Enough to evaluate in a single point θ_0 .
- ▶ **Kernel density estimator** to approximate $p(\theta_0|y)$. Unstable.
- ▶ Chib (1995) provide better solutions for **Gibbs sampling**.
- ▶ Chib and Jeliazkov (2001) generalizes to **MH algorithm** (good for Independence MH, not so good for RWM).

Computing the marginal likelihood: Approximation

- ▶ By **normal approximation** of the posterior distribution (**Lecture 6**).
- ▶ **Recall**: for large n

$$\begin{aligned} p(\theta|y) &\approx \mathcal{N}_p(\theta^*, \Sigma_{\theta^*} = J_{\theta^*,y}) \\ &= (2\pi)^{-p/2} |J_{\theta^*,y}|^{-1/2} \exp\left(-\frac{1}{2}(\theta - \theta^*)' J_{\theta^*,y}^{-1} (\theta - \theta^*)\right). \end{aligned}$$

- ▶ **The Laplace approximation**: Use rearranged Bayes' theorem with $\theta = \theta^*$

$$\log \hat{p}(y) = \log p(y|\theta^*) + \log p(\theta^*) + \frac{p}{2} \log(2\pi) + \frac{1}{2} \log |J_{\theta^*,y}^{-1}|.$$

- ▶ **As usual**: θ^* and $J_{\theta^*,y} [-H_{\theta^*}^{-1}]$ are obtained via a numerical optimization (e.g. `optim` in R).

Bayesian model averaging

- ▶ Let γ have the **same interpretation** across the same across the model space

$$\mathcal{M} = \{M_1, \dots, M_K\}.$$

Let $\theta = \{\theta_1, \dots, \theta_K\}$ be the corresponding set of parameters.

- ▶ The **marginal posterior** (marginalized over \mathcal{M}) of γ

$$p(\gamma|y) = \int p(\gamma, \mathcal{M}|y) d\mathcal{M} = \sum_{k=1}^K p(\gamma|M_k, y) p(M_k|y),$$

where $p(\gamma|M_k, y)$ is the **marginal posterior** (marginalized over θ_k) of γ conditional on model k ,

$$p(\gamma|M_k, y) = \int p(\gamma|\theta_k, y) p(\theta_k|y) d\theta_k.$$

- ▶ Note the **two layers** of averaging... **Bayes is all about averaging out (marginalize) unknown quantities!**

- ▶ **Example:** h -step ahead prediction for time series: $\gamma = (y_{T+1}, \dots, y_{T+h})$,

$$p(\gamma|M_k, y) = p_k(y_{T+1}, \dots, y_{T+h}|y) \quad [\text{Posterior predictive for } M_k]$$

$$p(M_k|y) \propto p(y|M_k)p(M_k), \quad [p(y|M_k) - \text{Marg. likelihood for } M_k]$$

- ▶ $p(y_{T+1}, \dots, y_{T+h}|y)$ includes **three sources of uncertainty**:

- ▶ **Future errors**/disturbances. **Simpler analogy:** σ^2 (assume known) in

$$y_1, \dots, y_n | \theta \stackrel{iid}{\sim} \mathcal{N}(\theta, \sigma^2), \quad \text{and } p(\theta) \propto c \text{ gives}$$

$$p(\theta|y) = \mathcal{N}(\bar{y}, \sigma^2/n)$$

$$p(\tilde{y}|y) = \mathcal{N}\left(\bar{y}, \sigma^2 + \frac{\sigma^2}{n}\right) \quad [\text{Posterior predictive for future } \tilde{y}].$$

- ▶ **Parameter uncertainty** (Posterior predictive averaged over posterior of θ).
- ▶ **Model uncertainty** (by model averaging).
- ▶ **Any painful integrals?** Compute by simulation!

Chib, S., (1995). Marginal likelihood from the Gibbs output. *Journal of the American Statistical Association*, 90(432):1313-1321

Chib, S. and Jeliazkov, I. (2001). Marginal likelihood from the MetropolisHastings output. *Journal of the American Statistical Association*, 96(453):270-281.

Lavine, M. and Schervish, M.J., (1999). Bayes factors: what they are and what they are not. *The American Statistician*, 53(2):119-122.